

Appendix for: "Multi-Aspect Conditioning for Diffusion-Based Music Synthesis: Enhancing Realism and Acoustic Control"

TABLE I
INSTRUMENT DISTRIBUTION OF THE TRAIN SET $\mathcal{D}_{\text{train}}$. THE TRAIN SET CONTAINS AUDIO AND CORRESPONDING MIDI. WE SHOW THE TOTAL LENGTH FOR EACH ENSEMBLE (LENGTH), AND HOW MANY PERFORMANCES (CORRESPONDING TO VERSIONS) ARE PLAYED BY EACH ENSEMBLE (#PERFORMANCES).

Ensemble	Length	#Performances
Flute & Harpsichord	0:49:19	2
Orchestra	14:31:13	39
Orchestra & Choir	1:49:54	1
Orchestra & Piano	5:12:50	7
Solo Cello	2:16:16	2
Solo Flute	0:09:16	1
Solo Guitar	4:40:14	85
Solo Harpsichord	7:27:40	12
Solo Organ	2:10:49	25
Solo Piano	5:52:41	17
Solo Violin	1:56:49	1
Violin & Harpsichord	3:18:48	2
Violin & Piano	3:59:17	1
Violin, Cello, & Piano	3:37:23	1
Wind Quintet	0:33:10	1
All	58:25:47	197

I. DATASET DISTRIBUTION

We provide the distribution of instrumentation in the different datasets used for training and evaluation, and a list of the pieces used for the listening tests. Table I shows the distribution of the train set $\mathcal{D}_{\text{train}}$. Table II shows the pieces appearing in the evaluation set used for the listening tests, which we refer to as $\mathcal{D}_{\text{listen}}$, together with their instrumentation. Table III shows the distribution of the large evaluation set used for quantitative evaluation, which we refer to as $\mathcal{D}_{\text{quant}}$.

II. QUANTITATIVE RESULTS OVER LARGE SCALE EVALUATION SET

In this section, we compare different models trained on the $\mathcal{D}_{\text{train}}$ dataset, appearing in Table I. We perform the evaluation on the large evaluation set $\mathcal{D}_{\text{quant}}$ appearing in Table III.

We compare the following models: U-Net and T5 models trained with alignments as score conditions (these two models appear in our previous work [1]), and a T5 model trained with automatic transcriptions as score conditions (evaluation in this paper was done using this model). For the latter model we provided two forms of score conditions: Pitch with instrument, and pitch-only. The transcriptions were obtained by training a transcriber [2, 3] on the pairs of audio and alignments, and providing the thresholded predictions as note conditions to the synthesizer, rather than the alignments.

Together, 4 different configurations were evaluated. We train all models both with, and without version conditioning – together 8 configurations.

Results are reported in Tables IV (All-FAD), V (Group-FAD), VI (version classification accuracy), and VII (transcription accuracy). The general trend we clearly and consistently observe is that Group-FAD and version classification accuracy dramatically improve as a result of version conditioning, while All-FAD and transcription metrics remain comparable (possibly with a slight increase or decrease). That is, by incorporating version conditioning, we can generate performances with the same notes, but such performances that more resemble the desired target versions, and do so while maintaining quality.

We also notice that training with transcription-based score conditions yields comparable results to alignment-based score conditions, with transcription-based conditions producing slightly better transcription scores, and alignment-based conditions producing slightly better FAD scores.

We describe the results in detail:

a) All-FAD: One can see in Table IV that version conditioning does not significantly impact the All-FAD – the VGGish All-FAD slightly increases while the TRILL All-FAD slightly decreases. Increase in All-FAD as a result of version conditioning is not necessarily surprising—it can be interpreted as follows: Version conditioning causes the generated performances to deviate from the general distribution of the the train set $\mathcal{D}_{\text{train}}$, towards the distribution of a subset \mathcal{D}_v corresponding to a specific version v .

The different compared models produce comparable results, except for the pitch-only model. It can be seen that conditioning on instrumentation improves the All-FAD score.

b) Group-FAD, Version Classification: The dramatic effect of version conditioning can be seen in Tables V and VI. The Group-FAD metric dramatically improves as a result of version conditioning, for all models, without exception. This means the distribution of the generated performances becomes perceptually more similar to the conditioning versions.

c) Transcription: We observe the transcription metrics (Table VII), measuring if the synthesized performances actually realize the notes specified by the MIDI note condition. There is no entirely objective or absolute way to measure this, however, we can still gain insights by using an automatic transcriber. We therefore use a transcriber trained on precisely the same data as the synthesizer. Note that such metrics are influenced not only by the quality of the synthesizer, but also from the quality of the transcriber.

In Table VII we can observe that most models yield similar transcription metrics, whether using version conditioning or

TABLE II

MIDI PERFORMANCES USED FOR THE LISTENING TEST WITH THEIR CORRESPONDING ENSEMBLES. THE FIRST THREE MINUTES OF EACH PERFORMANCE WERE USED. THIS DATASET IS REFERRED TO IN THE PAPER AS $\mathcal{D}_{\text{listen}}$. THIS DATASET CONTAINS BOTH MIDI AND CORRESPONDING AUDIO RECORDINGS OF REAL MUSICAL PERFORMANCES, OF THE EXACT SAME ENSEMBLES, OF VERSIONS THAT DO NOT APPEAR IN THE TRAIN SET.

MIDI	Ensemble
Bach Double Concerto in C Minor BWV 1060 (movement 1)	violin, cello, & harpsichord
Bach Great Fugue in G Minor BWV 542	harpsichord
Bach Italian Concerto in F Major BWV 971 (movement 1)	harpsichord
Bach Mass in B Minor BWV 232, Gloria - Cum Sancto Spiritu (movement 11)	choir & orchestra
Bach Orchestral Suite 1 in C Major BWV 1066, Overture	orchestra
Bach Orchestral Suite 2 in B Minor BWV 1067, Badinerie	wind quintet
Bach Toccata and Fugue in D Minor BWV 565	church organ
Beethoven Symphony 5 in C Minor Op. 67 (movement 1)	orchestra
Beethoven Symphony 6 in F Major Op. 68 (movement 1)	orchestra
Beethoven Symphony 6 in F Major Op. 68 (movement 3)	orchestra
Mozart Piano Concerto 20 in D Minor K. 466 (movement 1)	piano & orchestra
Mozart Symphony 40 in G Minor K. 550 (movement 1)	orchestra

TABLE III

INSTRUMENT DISTRIBUTION OF THE LARGE EVALUATION SET $\mathcal{D}_{\text{quant}}$, USED FOR QUANTITATIVE EVALUATION. THIS DATASET CONTAINS MIDI ONLY, WITHOUT CORRESPONDING AUDIO. WE SHOW THE TOTAL LENGTH FOR EACH ENSEMBLE (LENGTH), AND HOW MANY MIDI PERFORMANCES ARE PLAYED BY EACH ENSEMBLE (#PERFORMANCES).

Ensemble	Length	#Performances
Flute & Harpsichord	0:04:11	1
Orchestra	2:16:02	19
Orchestra & Choir	0:11:49	4
Orchestra & Piano	0:52:27	3
Solo Cello	0:05:53	3
Solo Guitar	0:05:07	3
Solo Harpsichord	0:09:17	4
Solo Organ	0:07:33	1
Solo Piano	0:18:53	2
Solo Violin	0:13:14	5
Violin & Harpsichord	0:05:44	2
Violin, Cello, & Piano	0:34:29	9
Wind Quintet	0:04:51	2
All	5:09:30	58

not, reaching accuracy of up to 67% (note-level), which is of reasonable magnitude when considering the complexity of highly polyphonic orchestral music. In addition, we can observe that using transcriptions as conditions rather than alignments provides better overall transcription metrics.

It can also be seen that unsurprisingly, the note-with-instrument metric is significantly lower when using pitch-only input. Note however that this is significantly mitigated when using version conditioning (15% improvement). This indicates that version conditioning helps achieving the target instrumentation, even when using pitch-only note conditions, as the model learns the correlations between versions and their instrumentation.

A. Difference in All-FAD vs. Group-FAD Value Range

When comparing the actual FAD values in Tables IV, V, one can see the All-FAD values are lower than the Group-FAD. We attribute this to the fact that the mean vectors and covariance matrices for All-FAD are computed over significantly larger evaluation and reference datasets than for Group-FAD and therefore yield less statistical fluctuations, resulting in an overall lower FAD score.

TABLE IV

ALL-FAD RESULTS ON LARGE EVALUATION SET $\mathcal{D}_{\text{quant}}$.

	All-FAD↓			
	VGGish		TRILL	
Version Cond.	w/o	w/	w/o	w/
U-Net Aligned	3.37	3.94	0.12	0.11
T5 Aligned	3.98	3.53	0.12	0.09
T5 Transcribed	3.05	3.58	0.12	0.11
T5 Transcribed Pitch	3.97	3.78	0.18	0.12

TABLE V

GROUP-FAD RESULTS ON LARGE EVALUATION SET $\mathcal{D}_{\text{quant}}$.

	Group-FAD↓			
	VGGish		TRILL	
Version Cond.	w/o	w/	w/o	w/
U-Net Aligned	7.06	5.21	0.5	0.33
T5 Aligned	6.95	5.46	0.51	0.35
T5 Transcribed	7.46	5.68	0.55	0.36
T5 Transcribed Pitch	12.81	6.03	0.7	0.38

TABLE VI

VERSION CLASSIFICATION RESULTS ON LARGE EVALUATION SET $\mathcal{D}_{\text{quant}}$.

Version Cond.	Classification% Top-1/3/5↑	
	w/o	w/
U-Net Aligned	11.6/24.5/36.1	52.9/73.5/84.5
T5 Aligned	35.5/60.0/69.7	67.7/89.7/91.0
T5 Transcribed	16.8/31.6/41.9	66.5/83.9/88.4
T5 Transcribed Pitch	4.5/18.1/27.1	62.6/79.4/85.8

TABLE VII

TRANSCRIPTION RESULTS ON LARGE EVALUATION SET $\mathcal{D}_{\text{quant}}$.

	Transcription F1%↑					
	Note		Note & Inst.		Frame	
Version Cond.	w/o	w/	w/o	w/	w/o	w/
U-Net Aligned	63.1	61.8	46.6	46.2	65.4	64.1
T5 Aligned	62.0	63.2	38.4	47.1	60.7	62.0
T5 Transcribed	66.9	64.7	50.7	46.2	64.8	63.9
T5 Transcribed Pitch	67.2	64.3	25.0	40.3	64.7	63.2

III. CLASSIFIER-FREE GUIDANCE WITH MULTIPLE CONDITIONS

Classifier-Free Guidance is a technique for controlling the condition strength in conditional diffusion models. The model

is trained both conditionally and unconditionally simultaneously, by applying condition dropout (zeroing out the condition), typically with probability 0.1. During sampling, noise is predicted both with and without the condition, and the enhanced conditioning is obtained through extrapolation in the condition’s ”direction”:

$$\begin{aligned}\epsilon_{cond} &= \epsilon_{\theta}(x_t, t, c) \\ \epsilon_{uncond} &= \epsilon_{\theta}(x_t, t, 0) \\ \epsilon &= \epsilon_{cond} + w(\epsilon_{cond} - \epsilon_{uncond})\end{aligned}\tag{1}$$

where $w > 0$ is an extrapolation weight controlling the desired conditioning strength. In our case of multi-aspect-conditioned music synthesis, we apply two simultaneous conditions: A score condition defining the notes to be played, and a version condition defining the acoustic- and performance-related properties such as timbre, recording environment, style, etc. It is possible to apply CFG using multiple conditions in a straightforward manner: Denote by c_1, c_2 the two types of conditions. During training, we dropout each of the conditions c_1, c_2 with probability 0.1, independently. Then, during sampling, we define:

$$\begin{aligned}\epsilon_{cond1} &= \epsilon_{\theta}(x_t, t, c1, 0) \\ \epsilon_{cond2} &= \epsilon_{\theta}(x_t, t, 0, c2) \\ \epsilon_{cond1,2} &= \epsilon_{\theta}(x_t, t, c1, c2) \\ \epsilon &= \epsilon_{cond1,2} + w_1(\epsilon_{cond1,2} - \epsilon_{cond2}) \\ &\quad + w_2(\epsilon_{cond1,2} - \epsilon_{cond1})\end{aligned}\tag{2}$$

where $w_1, w_2 > 0$ are extrapolation weights representing the strengths of the two conditions, respectively. Intuitively, the term $w_1(\epsilon_{cond1,2} - \epsilon_{cond2})$ enhances the condition c_1 while remaining faithful to the condition c_2 , and the term $w_2(\epsilon_{cond1,2} - \epsilon_{cond1})$ enhances the condition c_2 while remaining faithful to the condition c_1 . We used extrapolation weights of $w_1 = w_2 = 1.25$ for score and version after a grid search with values 1.0, 1.25, 1.5, 2.0 for w_1 and w_2 . We note that other ways to perform the extrapolation are possible, such as using $\epsilon_{\theta}(x_t, t, 0, 0)$ (zeroing out both conditions), but the approach described in Equation 2 gave best results in practice.

REFERENCES

- [1] B. Maman, J. Zeitler, M. Müller, and A. H. Bermanno, ”Performance conditioning for diffusion-based multi-instrument music synthesis,” in *Proceedings of the IEEE International Conference on Acoustics, Speech, and Signal Processing (ICASSP)*, Seoul, South Korea, 2024, pp. 5045–5049.
- [2] C. Hawthorne, A. Stasyuk, A. Roberts, I. Simon, C. A. Huang, S. Dieleman, E. Elsen, J. H. Engel, and D. Eck, ”Enabling factorized piano music modeling and generation with the MAESTRO dataset,” in *Proceedings of the International Conference on Learning Representations (ICLR)*, New Orleans, Louisiana, USA, 2019. [Online]. Available: <https://openreview.net/forum?id=r1lYRjC9F7>
- [3] B. Maman and A. H. Bermanno, ”Unaligned supervision for automatic music transcription in the wild,” in *Proceedings of the International Conference on Machine Learning (ICML)*, Baltimore, Maryland, USA, 2022, pp. 14 918–14 934.